

```
In [1]: pip install matplotlib
```

```
Requirement already satisfied: matplotlib in ./conda/envs/my5920_env/lib/python3.10/site-packages (3.10.0)
Requirement already satisfied: contourpy>=1.0.1 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from matplotlib) (1.3.1)
Requirement already satisfied: cyclor>=0.10 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from matplotlib) (4.56.0)
Requirement already satisfied: kiwisolver>=1.3.1 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from matplotlib) (1.4.8)
Requirement already satisfied: numpy>=1.23 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from matplotlib) (2.2.3)
Requirement already satisfied: packaging>=20.0 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from matplotlib) (24.2)
Requirement already satisfied: pillow>=8 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from matplotlib) (11.1.0)
Requirement already satisfied: pyparsing>=2.3.1 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from matplotlib) (3.2.1)
Requirement already satisfied: python-dateutil>=2.7 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from matplotlib) (2.9.0.post0)
Requirement already satisfied: six>=1.5 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)
Note: you may need to restart the kernel to use updated packages.
```

```
In [ ]: #problem 1.1
```

```
In [2]: import torch
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import Subset

# Define the transform to normalize
transform = transforms.Compose([
    transforms.ToTensor(), #image to tensor
    transforms.Normalize((0.1307,), (0.3081,)) # Normalize with mean and std of MN
])

# Loading the MNIST dataset
mnist_dataset = torchvision.datasets.MNIST(root='./data', train=True, download=True)

# Create a subset of 1,000 samples
subset_indices = torch.randperm(len(mnist_dataset))[:1000] # Randomly select 1,000
mnist_subset = Subset(mnist_dataset, subset_indices)

# Checking sample shape and label
sample_image, sample_label = mnist_subset[0]
print(f"Image Shape: {sample_image.shape}, Label: {sample_label}")
```

```
Image Shape: torch.Size([1, 28, 28]), Label: 8
```

```
In [16]: pip install pillow
```

Requirement already satisfied: pillow in ./conda/envs/my5920_env/lib/python3.10/site-packages (11.1.0)

Note: you may need to restart the kernel to use updated packages.


In []:

```
In [3]: import matplotlib.pyplot as plt
import torchvision.transforms as transforms

# a sample image from the MNIST subset
image, label = mnist_subset[1] # First image from the subset

# Convert from tensor (C, H, W) to (H, W) and plot
plt.imshow(image.squeeze(), cmap='gray') # Remove channel dimension and display
plt.title(f"Label: {label}")
plt.axis("off")
plt.show()
```

Label: 0




In []: *#problem 1.2*

```
In [4]: import matplotlib.pyplot as plt
# Extract labels from the subset
labels = [mnist_subset[i][1] for i in range(len(mnist_subset))]

# Plot histogram of class distribution
plt.figure(figsize=(8, 6))
plt.hist(labels, bins=range(11), align='left', rwidth=0.8, color='blue', edgecolor=
plt.xticks(range(10)) # Digits 0-9
plt.xlabel("Digit Class")
plt.ylabel("Frequency")
```

```
plt.title("Histogram of MNIST Subset Classes")
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```



In []: *#problem 1.3*

In [25]: `pip install einops`

Collecting einops

Downloading einops-0.8.1-py3-none-any.whl.metadata (13 kB)

Downloading einops-0.8.1-py3-none-any.whl (64 kB)

Installing collected packages: einops

Successfully installed einops-0.8.1

Note: you may need to restart the kernel to use updated packages.

```
In [5]: from torch.utils.data import Subset, DataLoader
        from einops import rearrange

        # Extract images and labels from the subset
        images = torch.stack([mnist_subset[i][0] for i in range(len(mnist_subset))])
        labels = torch.tensor([mnist_subset[i][1] for i in range(len(mnist_subset))])

        # Reshape using einops: [1000, 1, 28, 28] → [40, 25, 1, 28, 28]
        batch_size = 25
        batched_images = rearrange(images, '(b s) c h w -> b s c h w', s=batch_size)
        batched_labels = rearrange(labels, '(b s) -> b s', s=batch_size)
```

```
In [6]: print("Batched Image Shape:", batched_images.shape)
        print("Batched Label Shape:", batched_labels.shape)
```

```
Batched Image Shape: torch.Size([40, 25, 1, 28, 28])
Batched Label Shape: torch.Size([40, 25])
```

```
In [ ]: #problem 1.4
```

```
In [7]: import torch
import torchvision
import torchvision.transforms as transforms
import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D

# Randomly select one image sample
random_idx = torch.randint(0, len(mnist_dataset), (1,)).item()
image, label = mnist_dataset[random_idx]

# Convert the image tensor to a NumPy array
image_np = image.squeeze().numpy() # Shape: (28, 28)

# Generate x, y grid coordinates
x = np.arange(28)
y = np.arange(28)
X, Y = np.meshgrid(x, y)


# Plot the 3D surface
fig = plt.figure(figsize=(8, 6))
ax = fig.add_subplot(111, projection='3d')

# Create the surface plot
ax.plot_surface(X, Y, image_np, cmap='viridis', edgecolor='k')

# Labels and title
ax.set_xlabel("X (Pixel Location)")
ax.set_ylabel("Y (Pixel Location)")
ax.set_zlabel("Pixel Intensity")
ax.set_title(f"3D Plot of MNIST Image (Label: {label})")

plt.show()
```

3D Plot of MNIST Image (Label: 7)



```
In [39]: pip install pandas
```

```
Collecting pandas
  Downloading pandas-2.2.3-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (89 kB)
Requirement already satisfied: numpy>=1.22.4 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from pandas) (2.2.3)
Requirement already satisfied: python-dateutil>=2.8.2 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from pandas) (2.9.0.post0)
Collecting pytz>=2020.1 (from pandas)
  Downloading pytz-2025.1-py2.py3-none-any.whl.metadata (22 kB)
Collecting tzdata>=2022.7 (from pandas)
  Downloading tzdata-2025.1-py2.py3-none-any.whl.metadata (1.4 kB)
Requirement already satisfied: six>=1.5 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)
Downloading pandas-2.2.3-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (13.1 MB)
13.1/13.1 MB 116.5 MB/s eta 0:00:00
Downloading pytz-2025.1-py2.py3-none-any.whl (507 kB)
Downloading tzdata-2025.1-py2.py3-none-any.whl (346 kB)
Installing collected packages: pytz, tzdata, pandas
Successfully installed pandas-2.2.3 pytz-2025.1 tzdata-2025.1
Note: you may need to restart the kernel to use updated packages.
```

```
In [ ]: #problem 2.2.1
```

```
In [10]: # Load the dataset (replace with the correct path to the file)
df = pd.read_csv('energydata_complete.csv')
```

```
# Display the first few rows to inspect the data structure
print(df.head())
```

```
# Convert the 'date' column to datetime type
df['date'] = pd.to_datetime(df['date'])
```

	date	Appliances	lights	T1	RH_1	T2	RH_2	\
0	2016-01-11 17:00:00	60	30	19.89	47.596667	19.2	44.790000	
1	2016-01-11 17:10:00	60	30	19.89	46.693333	19.2	44.722500	
2	2016-01-11 17:20:00	50	30	19.89	46.300000	19.2	44.626667	
3	2016-01-11 17:30:00	50	40	19.89	46.066667	19.2	44.590000	
4	2016-01-11 17:40:00	60	40	19.89	46.333333	19.2	44.530000	

	T3	RH_3	T4	...	T9	RH_9	T_out	Press_mm_hg	\
0	19.79	44.730000	19.000000	...	17.033333	45.53	6.600000	733.5	
1	19.79	44.790000	19.000000	...	17.066667	45.56	6.483333	733.6	
2	19.79	44.933333	18.926667	...	17.000000	45.50	6.366667	733.7	
3	19.79	45.000000	18.890000	...	17.000000	45.40	6.250000	733.8	
4	19.79	45.000000	18.890000	...	17.000000	45.40	6.133333	733.9	

	RH_out	Windspeed	Visibility	Tdewpoint	rv1	rv2
0	92.0	7.000000	63.000000	5.3	13.275433	13.275433
1	92.0	6.666667	59.166667	5.2	18.606195	18.606195
2	92.0	6.333333	55.333333	5.1	28.642668	28.642668
3	92.0	6.000000	51.500000	5.0	45.410389	45.410389
4	92.0	5.666667	47.666667	4.9	10.084097	10.084097

[5 rows x 29 columns]

```
In [8]: import pandas as pd
import matplotlib.pyplot as plt
```

```
# Load the dataset (replace with the correct path to the file)
df = pd.read_csv('energydata_complete.csv')
```

```
# Display the first few rows to inspect the data structure
print(df.head())
```

```
# Convert the 'date' column to datetime type
df['date'] = pd.to_datetime(df['date'])
```

```
# Plotting energy consumption over the entire period
```

```
plt.figure(figsize=(12, 6))
plt.plot(df['date'], df['Appliances'], label='Appliance Energy Consumption', color=
plt.xlabel('Date')
plt.ylabel('Energy Consumption (Wh)')
plt.title('Appliance Energy Consumption Over Time')
plt.grid(True)
plt.legend()
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

```

# Plotting energy consumption for one week (let's choose the first week)
start_date = df['date'].min()
end_date = start_date + pd.Timedelta(weeks=1)

# Filter the data for this one-week period
df_one_week = df[(df['date'] >= start_date) & (df['date'] < end_date)]

# Plot the energy consumption for this week
plt.figure(figsize=(12, 6))
plt.plot(df_one_week['date'], df_one_week['Appliances'], label='Appliance Energy Co
plt.xlabel('Date')
plt.ylabel('Energy Consumption (Wh)')
plt.title('Appliance Energy Consumption for One Week')
plt.grid(True)
plt.legend()
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()


```

	date	Appliances	lights	T1	RH_1	T2	RH_2	\
0	2016-01-11 17:00:00	60	30	19.89	47.596667	19.2	44.790000	
1	2016-01-11 17:10:00	60	30	19.89	46.693333	19.2	44.722500	
2	2016-01-11 17:20:00	50	30	19.89	46.300000	19.2	44.626667	
3	2016-01-11 17:30:00	50	40	19.89	46.066667	19.2	44.590000	
4	2016-01-11 17:40:00	60	40	19.89	46.333333	19.2	44.530000	

	T3	RH_3	T4	...	T9	RH_9	T_out	Press_mm_hg	\
0	19.79	44.730000	19.000000	...	17.033333	45.53	6.600000	733.5	
1	19.79	44.790000	19.000000	...	17.066667	45.56	6.483333	733.6	
2	19.79	44.933333	18.926667	...	17.000000	45.50	6.366667	733.7	
3	19.79	45.000000	18.890000	...	17.000000	45.40	6.250000	733.8	
4	19.79	45.000000	18.890000	...	17.000000	45.40	6.133333	733.9	

	RH_out	Windspeed	Visibility	Tdewpoint	rv1	rv2
0	92.0	7.000000	63.000000	5.3	13.275433	13.275433
1	92.0	6.666667	59.166667	5.2	18.606195	18.606195
2	92.0	6.333333	55.333333	5.1	28.642668	28.642668
3	92.0	6.000000	51.500000	5.0	45.410389	45.410389
4	92.0	5.666667	47.666667	4.9	10.084097	10.084097

[5 rows x 29 columns]



```
In [ ]: #problem 2.2.2
```

```
In [16]: import seaborn as sns

# Select the data for one week (e.g., the first 7 days of data)
start_date = df['date'].min()
end_date = start_date + pd.Timedelta(weeks=1)

# Filter data for the one-week period
df_one_week = df[(df['date'] >= start_date) & (df['date'] < end_date)]

# Set the 'date' column as the index for resampling
df_one_week.set_index('date', inplace=True)

# Resample the data to hourly averages (since the data is in 10-minute intervals)
df_hourly = df_one_week['Appliances'].resample('H').mean()


# Create a new DataFrame for the heatmap, which will be structured by hour of the d
df_heatmap = df_hourly.groupby([df_hourly.index.hour, df_hourly.index.weekday]).mea
```



```
# Plotting the heatmap
plt.figure(figsize=(12, 6))
sns.heatmap(df_heatmap, cmap='YlOrRd', annot=True, fmt='.1f', cbar_kws={'label': 'E
plt.ylabel('Hour of Day')
plt.xlabel('Day of Week')
plt.title('Heatmap of Hourly Appliance Energy Consumption for One Week')
plt.yticks(ticks=range(0, 24), labels=[f'{i}:00' for i in range(24)], rotation=45)
plt.xticks(ticks=range(7), labels=['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun'])
plt.tight_layout()
plt.show()
```

/tmp/ipykernel_513178/345016752.py:14: FutureWarning: 'H' is deprecated and will be removed in a future version, please use 'h' instead.

```
df_hourly = df_one_week['Appliances'].resample('H').mean()
```




In []: `#problem 2.2.3`

```
In [45]: # Plotting the histogram of energy consumption
plt.figure(figsize=(10, 6))
sns.histplot(df['Appliances'], bins=30, kde=True, color='skyblue', edgecolor='black')

# Adding titles and labels
plt.title('Histogram of Energy Consumption of Appliances', fontsize=16)
plt.xlabel('Energy Consumption (Wh)', fontsize=14)
plt.ylabel('Frequency', fontsize=14)

# Show the plot
plt.tight_layout()
plt.show()
```



```
In [ ]: #problem 2.2.4
```

```
In [46]: # Convert 'date' to datetime
df['date'] = pd.to_datetime(df['date'])


# Extract the time part from the 'date' column
df['time'] = df['date'].dt.time

# Calculate the number of seconds from midnight for each entry
df['NSM'] = df['date'].dt.hour * 3600 + df['date'].dt.minute * 60 + df['date'].dt.second

# Plotting energy consumption vs. NSM (number of seconds from midnight)
plt.figure(figsize=(10, 6))
plt.scatter(df['NSM'], df['Appliances'], alpha=0.5, color='blue')

# Adding titles and labels
plt.title('Energy Consumption vs. Number of Seconds from Midnight (NSM)', fontsize=14)
plt.xlabel('Number of Seconds from Midnight (NSM)', fontsize=14)
plt.ylabel('Energy Consumption (Wh)', fontsize=14)

# Show the plot
plt.tight_layout()
plt.show()
```

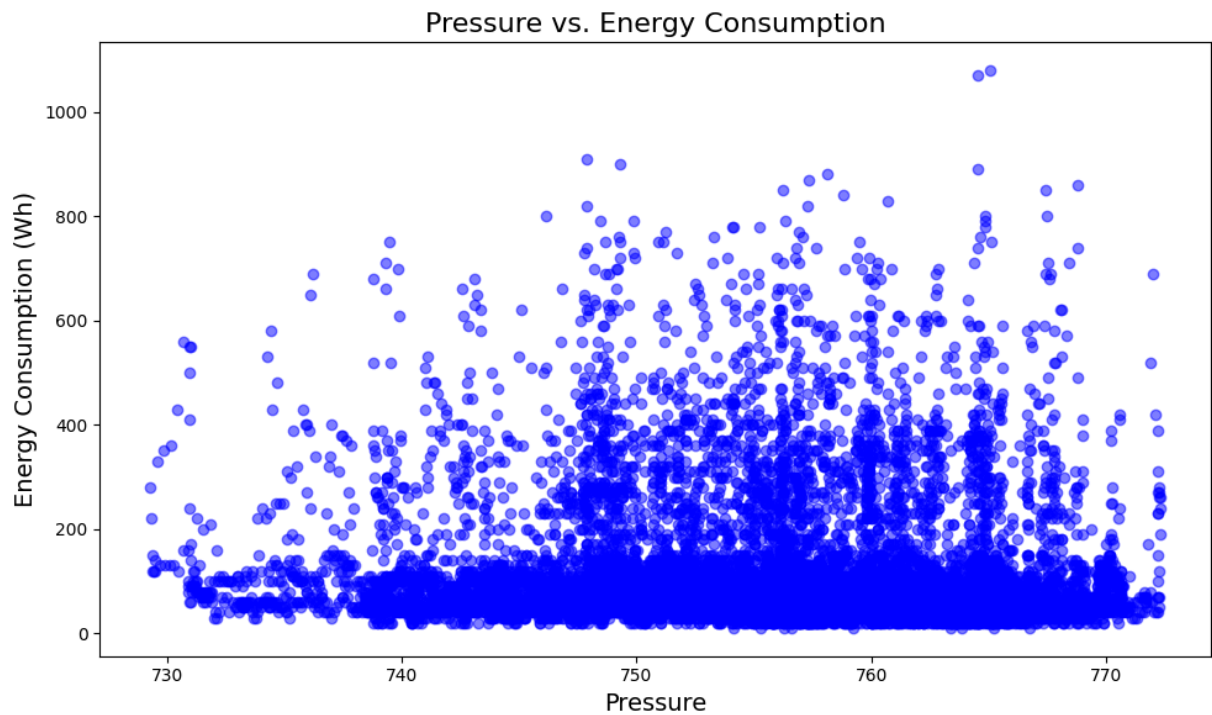


In []: *#problem 2.2.5*

```
In [47]: plt.figure(figsize=(10, 6))
plt.scatter(df['Press_mm_hg'], df['Appliances'], alpha=0.5, color='blue')

# Adding titles and labels
plt.title('Pressure vs. Energy Consumption', fontsize=16)
plt.xlabel('Pressure', fontsize=14)
plt.ylabel('Energy Consumption (Wh)', fontsize=14)

# Show the plot
plt.tight_layout()
plt.show()
```



```
In [ ]: ## problem 2.2.6
```

```
##Press mm Hg (pressure) and NSM (time of day) are the main determinants of energy
##Press mm Hg controls energy needs through atmospheric conditions, whereas NSM inf
```

```
In [2]: pip install scipy
```

Collecting scipy

Downloading scipy-1.15.2-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (61 kB)

Requirement already satisfied: numpy<2.5,>=1.23.5 in ./conda/envs/my5920_env/lib/python3.10/site-packages (from scipy) (2.2.3)

Downloading scipy-1.15.2-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (37.6 MB)

37.6/37.6 MB 140.8 MB/s eta 0:00:0000:01

Installing collected packages: scipy

Successfully installed scipy-1.15.2

Note: you may need to restart the kernel to use updated packages.

```
In [ ]: #Problem 3
```

```
In [3]: import pandas as pd
import numpy as np
from scipy.stats import kurtosis, skew
```

```
# Load the data (modify delimiter based on file structure)
```

```
file_path = "airfoil_self_noise.dat" # Replace with actual file path
```

```
data = pd.read_csv(file_path, delim_whitespace=True, header=None)
```

```
# Assign column names (assuming there are six columns based on the description)
```

```
column_names = [
```

```
    "Frequency (Hz)",
```

```
    "Angle of Attack (°)",
```

```

    "Chord Length (m)",
    "Free-stream Velocity (m/s)",
    "Suction Side Displacement Thickness (m)",
    "Scaled Sound Pressure Level (dB)"
]
data.columns = column_names

# Compute descriptive statistics
stats = {}
for column in data.columns:
    stats[column] = {
        "Mean": np.mean(data[column]),
        "Variance": np.var(data[column], ddof=1),
        "Standard Deviation": np.std(data[column], ddof=1),
        "Median": np.median(data[column]),
        "Kurtosis": kurtosis(data[column]),
        "Skewness": skew(data[column]),
        "Range": np.ptp(data[column])
    }

# Convert to DataFrame for better visualization
stats_df = pd.DataFrame(stats).T
print(stats_df)

```

	Mean	Variance \
Frequency (Hz)	2886.380572	9.938717e+06
Angle of Attack (°)	6.782302	3.502424e+01
Chord Length (m)	0.136548	8.749868e-03
Free-stream Velocity (m/s)	50.860745	2.425116e+02
Suction Side Displacement Thickness (m)	0.011140	1.729287e-04
Scaled Sound Pressure Level (dB)	124.835943	4.759146e+01

	Standard Deviation	Median \
Frequency (Hz)	3152.573137	1600.000000
Angle of Attack (°)	5.918128	5.400000
Chord Length (m)	0.093541	0.101600
Free-stream Velocity (m/s)	15.572784	39.600000
Suction Side Displacement Thickness (m)	0.013150	0.004957
Scaled Sound Pressure Level (dB)	6.898657	125.721000

	Kurtosis	Skewness	Range
Frequency (Hz)	5.685722	2.134951	19800.000000
Angle of Attack (°)	-0.415568	0.688476	22.200000
Chord Length (m)	-1.038504	0.457001	0.279400
Free-stream Velocity (m/s)	-1.562743	0.235617	39.600000
Suction Side Displacement Thickness (m)	2.207539	1.700465	0.058011
Scaled Sound Pressure Level (dB)	-0.317132	-0.418534	37.607000

/tmp/ipykernel_488093/32569673.py:7: FutureWarning: The 'delim_whitespace' keyword in pd.read_csv is deprecated and will be removed in a future version. Use ``sep='\s+'`` instead

```
data = pd.read_csv(file_path, delim_whitespace=True, header=None)
```